

Optimization Algorithms from Scratch Project Report

Contributors

- **Sumit Hulke** (2023BCD0026)
 - **Suraj Sanjay Harlekar** (2023BCD0038)
 - **Aryan Patil** (2023BCD0047)
-

1. Introduction

This project focuses on implementing a complete suite of **optimization algorithms from scratch** using only **NumPy**.

Our goal is to understand:

- How different optimization algorithms work internally
- How they behave on classical test functions
- How fast they converge
- How their convergence trajectories differ when visualized on contour plots

The project is structured modularly, with algorithms defined in Python modules and experiments performed using Jupyter Notebooks.



2. Algorithms Implemented

Below is a brief explanation of every algorithm covered in the project.

2.1 First-Order Optimization Methods

Gradient Descent (GD)

- Moves opposite to the gradient direction.
- Highly dependent on learning rate.



Stochastic Gradient Descent (SGD)

- Uses one randomly sampled data point.

- Noisy updates → helps escape local minima.

Mini-Batch SGD

- Uses small batches for smoother, faster convergence.
-


2.2 Momentum-Based Methods

Momentum Gradient Descent

- Adds a velocity term to gradient updates.
- Reduces oscillations, especially in curved valleys.

Nesterov Accelerated Gradient (NAG)

- Computes gradient at a look-ahead position.
- Results in faster and more stable convergence.



2.3 Adaptive Learning Rate Methods

Adagrad


- Adapts learning rates per-parameter based on history.
- Great for sparse features.

RMSProp

- Uses moving average of squared gradients.
- Prevents learning rate from shrinking too fast.

Adam

- Combines Momentum + RMSProp.
- Default choice in deep learning.




2.4 Second-Order Optimization Methods

Newton's Method

- Uses Hessian matrix to take curvature-aware steps.
- Extremely fast near minima.

Damped Newton Method

- Adds damping factor when the Hessian is not positive definite.



2.5 Quasi-Newton Methods



BFGS


- Approximates the inverse Hessian using gradient differences.

- Faster than vanilla GD, without requiring true Hessians.

L-BFGS

- Limited memory BFGS (stores only a few past updates).
- Suitable for high-dimensional problems.






2.6 Non-Differentiable Optimization

Sub-Gradient Method

- Works when objective is not differentiable (e.g., L1 norm).
- Uses any valid sub-gradient at non-smooth points.



3. Regression Models Implemented

We also implemented regression models (from scratch + using optimizers).

3.1 Linear Regression

- Uses closed-form solution and GD.
- Dataset: **California Housing**.

3.2 Ridge Regression (L2 Regularization)



- Adds L2 penalty to control model complexity.


3.3 Lasso Regression (L1 Regularization)

- Uses subgradient method.
- Produces sparse parameter vectors.

3.4 Logistic Regression

- Binary classification using sigmoid + cross-entropy loss.
- Dataset: **Breast Cancer Wisconsin**.





4. Project Structure Summary

- **optimizers/** → All algorithm implementations
 - **notebooks/** → Experiments, contour plots, convergence graphs
 - **utils/** → Test functions, plotting helpers
-

5. Conclusion

This project provided deep insight into:

- Gradient-based optimization
- Curvature-aware second-order methods
- Adaptive learning rate strategies
- Quasi-Newton approximations

- Regression model training from scratch